

SQL Basics

What is it?

- Stands for “Structured Query Language”
- Standard for storing and accessing information
- 3 Major “sub-languages” or components to SQL
 - DDL - Data Declaration Language
 - DML - Data Manipulation Language
 - DCL - Data Control Language

DDL Commands

- Commands affect the structure and storage of data
- Create new tables:
`CREATE TABLE <table name> (<columns>);`
- Delete tables that already exist
`DROP TABLE <table name>;`
- Change the table definition:
`ALTER TABLE <table name> ... ;`

DML - Data Manipulation Language

- Commands access and modify to the *contents* of database
- Read values out of the database:
`SELECT * FROM users;`
- Put values into tables:
`INSERT INTO users ...;`
- Remove values from the database
`DELETE FROM users ...;`

DCL - Data Control Language

- Commands operate on data permissions
- Allow users to read data from a table:
`GRANT SELECT ON <table> TO <user>;`
- Prohibit users from modifying specific tables:
`REVOKE UPDATE ON <table> FROM <user>;`
- Won't go into more depth. Differs between database implementations significantly.

SQL is **old**. Why are we still talking about it?

- Declarative language that describes “what you want” not “how to do it”
- Friendly format for both humans and machines
- Lack of better general-purpose alternatives
- Extremely flexible abstractions

Basic Abstractions

- A SQL database stores data in one more more tables
- Tables are a collection of columns that have names and data types
- Rows are entries in tables that may or may not have values for each column
- Primary Keys uniquely identify rows within the table

id	email_address	first_name	last_name	is_admin
1	<u>jpetty@harrys.com</u>	James	Petty	TRUE
2	<u>chris@harrys.com</u>	Chris	Clouten	TRUE
3	<u>bigfoot@gmail.com</u>	NULL	NULL	FALSE
4	<u>imissthe70s@aol.com</u>	Cher	NULL	FALSE

- Primary Key?
- Columns?
- Rows?
- What's this NULL thing?

DDL - Data Definition Language

```
CREATE TABLE users (  
  id integer primary key,  
  email_address varchar not null,  
  first_name varchar,  
  last_name varchar,  
  is_admin boolean not null default false  
);
```

id	email_address	first_name	last_name	is_admin
1	jpetty@harrys.com	James	Petty	TRUE
2	chris@harrys.com	Chris	Clouten	TRUE
3	bigfoot@gmail.com	NULL	NULL	FALSE
4	imissthe70s@aol.com	Cher	NULL	FALSE

Common SQL Data Types

Name	Description	Values
boolean	Values of either “yes” or “no”	{true, false}
integer	Whole numbers (including 0) including negative values.	{..., -2, -1, 0, 1, 2, ...}
decimal(p,s)	Decimal value with precision p (total number of digits) and scale (number of fractional digits)	eg: decimal(5,2): [-999.99, 999.99]
char(n)	Text value with exact length (n)	Σ^n
varchar(n)	Text value with a variable length up to n .	$\Sigma^* = \bigcup \Sigma^n$ $n \in \mathbb{N} \cup \{0\}$

More SQL Data Types

Name	Description	Values
date	A date value including year, month, and day.	Varies by implementation
time	A time of day including hour, minute, second, and sometimes fractional seconds	00:00:00 - 24:00:00
timestamp	Combination of date with time, sometimes with timezone	Varies by implementation

DDL Column Definitions

- Columns *must* have a name and data type
- Columns *may* have constraints and a default value

name	data type	constraints	default value
id	integer	primary key	(it's complicated)
email_address	varchar	not null	None
first_name	varchar	None	NULL
last_name	varchar	None	NULL
is_admin	boolean	not null	FALSE

Interactive Time

- sqlite3 is a small, embeddable, SQL database engine that comes preinstalled with OS X
- make sure we can all run commands
- materials located here:

<https://pettyjamesm.github.io/mammoth-school-sql-intro/>

SQL Basics

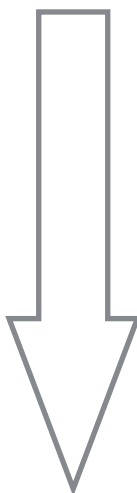
Part II, The Sequel

Recap

- SQL stands for “Structured Query Language.”
 - 3 major “sub-languages” or components to SQL
 - DDL - Data Declaration Language
 - DML - Data Manipulation Language
 - DCL - Data Control Language
 - DML is used to create, read, update, and delete information.

Tables / Relations

Columns



id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	TRUE
2	chris@example.com	Chris	Clouten	TRUE
3	bigfoot@gmail.com	NULL	NULL	FALSE
4	imissthe70s@aol.com	Cher	NULL	FALSE

Rows

- 
- Types (varchar, integer, boolean, ...)
 - “No Value” marker NULL

DML — INSERT

users

id	email_address	first_name	last_name	is_admin
Keys				

INSERT INTO users

(id, email_address, first_name, last_name, is_admin)

VALUES

(1, 'jpetty@example.com', 'James', 'Petty', 1),

(2, 'chris@example.com', 'Chris', 'Clouten', 1);

Result

users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1 (true)
2	chris@example.com	Chris	Clouten	1 (true)

Remember the **DDL**...

```
CREATE TABLE users (  
  id integer primary key,  
  email_address varchar not null,  
  first_name varchar,  
  last_name varchar,  
  is_admin boolean not null default false  
);
```

Constraints

```
INSERT INTO users (is_admin) VALUES (0);
```

Error: NOT NULL constraint failed: users.email_address

```
INSERT INTO users (email_address, is_admin)  
VALUES ('bigfoot@gmail.com', 0);
```

```
INSERT INTO users (id, email_address, is_admin)  
VALUES (3, 'imissthe70s@aol.com', 0);
```

Error: UNIQUE constraint failed: users.id

Constraints

- Database is checking consistency
 - Value in the correct domain?
 - Value specified?
 - Value unique?
 - Default?
- And more.

DML — SELECT

users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1
2	chris@example.com	Chris	Clouten	1
3	bigfoot@gmail.com	NULL	NULL	0

```
SELECT id, email_address  
FROM users;
```

```
1 | jpetty@example.com  
2 | chris@example.com  
3 | bigfoot@gmail.com
```

SELECT ... WHERE

users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1
2	chris@example.com	Chris	Clouten	1
3	bigfoot@gmail.com	NULL	NULL	0

```
SELECT id, email_address
FROM users
WHERE last_name IS NULL;
```

3 | bigfoot@gmail.com

count(), max(), etc.

users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1 (true)
2	chris@example.com	Chris	Clouten	1 (true)
3	bigfoot@gmail.com	NULL	NULL	0 (false)

```
SELECT count(*)  
FROM users  
WHERE is_admin;
```

2

```
SELECT max(id)  
FROM users;
```

3

SELECT * ... LIMIT

users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1
2	chris@example.com	Chris	Clouten	1
3	bigfoot@gmail.com	NULL	NULL	0

```
SELECT *  
FROM users  
LIMIT 1;
```

2 | chris@example.com | Chris | Clouten | 1

SELECT ... ORDER BY

 products

id	name	price
1	Daily Face Wash	7.00
2	Truman Set	15.00
3	Razor Blades	16.00
4	Foaming Shave Gel	6.00

```
SELECT name, price FROM products
ORDER BY price;
```

Foaming Shave Gel | 6
Daily Face Wash | 7
Truman Set | 15
Razor Blades | 16


JOINS

- Tables may be “joined” conceptually combining every row from one table with every row from the other table.
- Usually, you want a subset of the resulting rows...
- Each table stores data to model part of your problem — pulling the data together allows queries to look across tables for the answers.

Orders ➡ Users

orders

id	user_id	placed_at
1	2	2015-12-04 06:02:30
2	4	2016-02-25 04:45:14
3	3	2016-02-24 04:12:05



users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1
2	chris@example.com	Chris	Clouten	1
3	bigfoot@gmail.com	NULL	NULL	0
4	imissthe70s@aol.com	Cher	NULL	0

Orders ➡ Users

orders

id	user_id	placed_at
1	2	2015-12-04 06:02:30
2	4	2016-02-25 04:45:14
3	3	2016-02-24 04:12:05

users

id	email_address	first_name	last_name	is_admin
1	jpetty@example.com	James	Petty	1
2	chris@example.com	Chris	Clouten	1
3	bigfoot@gmail.com	NULL	NULL	0
4	imissthe70s@aol.com	Cher	NULL	0

id	email_address	first_name	last_name	is_admin
4	imissthe70s@aol.com	Cher	NULL	0

id	email_address	first_name	last_name	is_admin
3	bigfoot@gmail.com	NULL	NULL	0



id	user_id	placed_at	id	email_address	first_name	last_name	is_admin
1	2	2015-12-04 06:02:30	2	chris@example.com	Chris	Clouten	1
2	4	2016-02-25 04:45:14	4	imissthe70s@aol.com	Cher	NULL	0
3	3	2016-02-24 04:12:05	3	bigfoot@gmail.com	NULL	NULL	0

```
SELECT orders.placed_at, users.email_address, users.first_name
FROM orders, users
WHERE orders.user_id = users.id;
```

```
2015-12-04 06:02:30 | chris@example.com | Chris
2016-02-25 04:45:14 | imissthe70s@aol.com | Cher
2016-02-24 04:12:05 | bigfoot@gmail.com | NULL
```

FROM... JOIN ... ON

Orders

Users

id	user_id	placed_at	id	email_address	first_name	last_name	is_admin
1	2	2015-12-04 06:02:30	2	chris@example.com	Chris	Clouten	1
2	4	2016-02-25 04:45:14	4	imissthe70s@aol.com	Cher	NULL	0
3	3	2016-02-24 04:12:05	3	bigfoot@gmail.com	NULL	NULL	0

```
SELECT orders.placed_at, users.email_address, users.first_name
FROM orders
JOIN users ON orders.user_id = users.id;
```

```
2015-12-04 06:02:30 | chris@example.com | Chris
2016-02-25 04:45:14 | imissthe70s@aol.com | Cher
2016-02-24 04:12:05 | bigfoot@gmail.com | NULL
```

Interactive Time

- We will use sqlite3 again.
- Start terminal (using Spotlight), then type: **sqlite3**
- Online materials located here:

<https://pettyjamesm.github.io/mammoth-school-sql-intro/>